

```
1 import numpy as np
2 import pandas as pd
3 import tensorflow as tf
4 import datetime
```

▼ Loading and running our previous MNIST handwritten digits dataset

- Using 20 Epochs

```
1 from keras.datasets import mnist
2 from keras.utils import to_categorical
3
4 (train_images, train_labels), (test_images, test_labels) = mnist.load_data()
5
6 train_images = train_images.reshape((60000, 28, 28, 1))
7 train_images = train_images.astype('float32') / 255
8
9 test_images = test_images.reshape((10000, 28, 28, 1))
10 test_images = test_images.astype('float32') / 255
11
12 train_labels = to_categorical(train_labels)
13 test_labels = to_categorical(test_labels)
14
15 from keras import layers
16 from keras import models
17 from keras.layers import Dropout
18
19 model = models.Sequential()
20 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
21 model.add(Dropout(0.5))
22 model.add(layers.MaxPooling2D((2, 2)))
23 model.add(layers.Conv2D(16, (3, 3), activation='relu'))
24 model.add(layers.Flatten())
25 model.add(layers.Dense(16, activation='relu'))
26 model.add(Dropout(0.5))
27 model.add(layers.Dense(10, activation='sigmoid'))
28
29 model.summary()
30
31 model.compile(optimizer='rmsprop',
32               loss='categorical_crossentropy',
33               metrics=['accuracy'])
34
35 model.fit(train_images, train_labels,
36           epochs=20,
37           batch_size=64,
38           validation_split = 0.2)
```

dropout (Dropout)	(None, 26, 26, 32)	0
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 16)	4624
flatten (Flatten)	(None, 1936)	0
dense (Dense)	(None, 16)	30992
dropout_1 (Dropout)	(None, 16)	0
dense_1 (Dense)	(None, 10)	170
=====		
Total params: 36,106		
Trainable params: 36,106		
Non-trainable params: 0		

Epoch 1/20

750/750 [=====] - 44s 57ms/step - loss: 1.3461 - accuracy: 0.0000

Epoch 2/20

750/750 [=====] - 42s 57ms/step - loss: 0.5766 - accuracy: 0.0000

Epoch 3/20

750/750 [=====] - 43s 57ms/step - loss: 0.4781 - accuracy: 0.0000

Epoch 4/20

750/750 [=====] - 43s 57ms/step - loss: 0.4366 - accuracy: 0.0000

Epoch 5/20

750/750 [=====] - 43s 58ms/step - loss: 0.4164 - accuracy: 0.0000

Epoch 6/20

750/750 [=====] - 44s 59ms/step - loss: 0.4172 - accuracy: 0.0000

Epoch 7/20

750/750 [=====] - 44s 59ms/step - loss: 0.4150 - accuracy: 0.0000

Epoch 8/20

750/750 [=====] - 44s 59ms/step - loss: 0.4064 - accuracy: 0.0000

Epoch 9/20

750/750 [=====] - 45s 60ms/step - loss: 0.4108 - accuracy: 0.0000

Epoch 10/20

750/750 [=====] - 45s 60ms/step - loss: 0.4087 - accuracy: 0.0000

Epoch 11/20

750/750 [=====] - 44s 59ms/step - loss: 0.4174 - accuracy: 0.0000

Epoch 12/20

750/750 [=====] - 45s 60ms/step - loss: 0.4185 - accuracy: 0.0000

Epoch 13/20

750/750 [=====] - 46s 61ms/step - loss: 0.4139 - accuracy: 0.0000

Epoch 14/20

750/750 [=====] - 42s 57ms/step - loss: 0.4172 - accuracy: 0.0000

Epoch 15/20

750/750 [=====] - 43s 58ms/step - loss: 0.4173 - accuracy: 0.0000

Epoch 16/20

750/750 [=====] - 43s 58ms/step - loss: 0.4192 - accuracy: 0.0000

Epoch 17/20

750/750 [=====] - 44s 59ms/step - loss: 0.4132 - accuracy: 0.0000

Epoch 18/20

750/750 [=====] - 46s 61ms/step - loss: 0.4186 - accuracy: 0.0000

Epoch 19/20

750/750 [=====] - 43s 57ms/step - loss: 0.4178 - accuracy: 0.0000

Epoch 20/20

```
750/750 [=====] - 44s 59ms/step - loss: 0.4227 - accuracy: 0.7515 - 0.7515
```

Evaluating the model on the test set

```
1 model.save_weights('model.h1')
2 model.load_weights('model.h1')
3 model.evaluate(test_images, test_labels)
```

```
313/313 [=====] - 2s 8ms/step - loss: 0.2068 - acc: 0.6001
[0.2068410962820053, 0.6001999974250793]
```

▼ Automatically reducing overall runtime by using earlystopping

- we can use `callback_early_stopping` to interrupt training once a target metric being monitored has stopped improving for a fixed number of epochs.
- this callback allows us to interrupt training as soon as we start overfitting, thus avoiding having to retrain our model for a smaller number of epochs.

MODEL CHECKPOINT AND EARLY STOPPING CALLBACKS

```
1 import keras
2
3 callbacks_list = [
4     keras.callbacks.EarlyStopping(
5         monitor='acc',
6         patience=1,
7     ),
8     keras.callbacks.ModelCheckpoint(
9         filepath='my_model.h5',
10        monitor='val_loss',
11        save_best_only=True,
12    )
13 ]
14
15 model.compile(optimizer='rmsprop',
16              loss='binary_crossentropy',
17              metrics=['acc'])
18
19 model.fit(train_images, train_labels,
20         epochs=20,
21         batch_size=64,
22         callbacks=callbacks_list,
23         validation_data=(test_images, test_labels))
```

Epoch 1/20

```
938/938 [=====] - 57s 60ms/step - loss: 0.0984 - acc: 0.9999
```

Epoch 2/20

938/938 [=====] - 57s 61ms/step - loss: 0.0951 - acc: 0
 <tensorflow.python.keras.callbacks.History at 0x7f159f3fc690>

REDUCING LEARNING RATE ON PLATEAU CALLBACK

```
1 callbacks_list = [
2     keras.callbacks.ReduceLROnPlateau(
3         monitor='val_loss',
4         factor=0.1,
5         patience=10,
6     )
7 ]
8
9 model.fit(train_images, train_labels,
10          epochs=20,
11          batch_size=32,
12          callbacks=callbacks_list,
13          validation_data=(test_images, test_labels))
```

```
Epoch 1/20
1875/1875 [=====] - 64s 34ms/step - loss: 0.0974 - acc:
Epoch 2/20
1875/1875 [=====] - 62s 33ms/step - loss: 0.0987 - acc:
Epoch 3/20
1875/1875 [=====] - 63s 33ms/step - loss: 0.0994 - acc:
Epoch 4/20
1875/1875 [=====] - 65s 35ms/step - loss: 0.0992 - acc:
Epoch 5/20
1875/1875 [=====] - 65s 35ms/step - loss: 0.0994 - acc:
Epoch 6/20
1875/1875 [=====] - 66s 35ms/step - loss: 0.0998 - acc:
Epoch 7/20
1875/1875 [=====] - 68s 36ms/step - loss: 0.1001 - acc:
Epoch 8/20
1875/1875 [=====] - 67s 36ms/step - loss: 0.0991 - acc:
Epoch 9/20
1875/1875 [=====] - 65s 35ms/step - loss: 0.1000 - acc:
Epoch 10/20
1875/1875 [=====] - 65s 34ms/step - loss: 0.1004 - acc:
Epoch 11/20
1875/1875 [=====] - 64s 34ms/step - loss: 0.1000 - acc:
Epoch 12/20
1875/1875 [=====] - 62s 33ms/step - loss: 0.0997 - acc:
Epoch 13/20
1875/1875 [=====] - 62s 33ms/step - loss: 0.0999 - acc:
Epoch 14/20
1875/1875 [=====] - 60s 32ms/step - loss: 0.0990 - acc:
Epoch 15/20
1875/1875 [=====] - 61s 33ms/step - loss: 0.0994 - acc:
Epoch 16/20
1875/1875 [=====] - 61s 33ms/step - loss: 0.0996 - acc:
Epoch 17/20
1875/1875 [=====] - 62s 33ms/step - loss: 0.1004 - acc:
Epoch 18/20
```

```
1875/1875 [=====] - 62s 33ms/step - loss: 0.0994 - acc:
Epoch 19/20
1875/1875 [=====] - 62s 33ms/step - loss: 0.1002 - acc:
Epoch 20/20
1875/1875 [=====] - 61s 33ms/step - loss: 0.1009 - acc:
<tensorflow.python.keras.callbacks.History at 0x7f15a26c9b90>
```

```
1 import keras
2 import numpy as np
3
4 class ActivationLogger(keras.callbacks.Callback):
5
6     def set_model(self, model):
7         self.model = model
8         layer_outputs = [layer.output for layer in model.layers]
9         self.activations_model = keras.models.Model(model.input,
10                                                       layer_outputs)
11
12     def on_epoch_end(self, epoch, logs=None):
13         if self.validation_data is None:
14             raise RuntimeError('Requires validation_data.')
15
16         validation_sample = self.validation_data[0][0:1]
17         activations = self.activations_model.predict(validation_sample)
18         f = open('activations_at_epoch_' + str(epoch) + '.npz', 'w')
19         np.savez(f, activations)
20         f.close()
```

Training the model with a TensorBoard callback

```
1 from keras.callbacks import TensorBoard
2 tbCallBack = TensorBoard(log_dir='./log', histogram_freq=1,
3                           write_graph=True,
4                           write_grads=True,
5                           batch_size=128,
6                           write_images=True)
7
8 model.fit(train_images, train_labels,
9           batch_size=128,
10           epochs=20,
11           verbose=1,
12           validation_data=(test_images, test_labels),
13           callbacks=[tbCallBack])
```

```
WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback
WARNING:tensorflow:`batch_size` is no longer needed in the `TensorBoard` Callback
Epoch 1/20
469/469 [=====] - 48s 102ms/step - loss: 0.0944 - acc: 0.9999
Epoch 2/20
469/469 [=====] - 48s 102ms/step - loss: 0.0941 - acc: 0.9999
```

```

Epoch 3/20
469/469 [=====] - 48s 102ms/step - loss: 0.0940 - acc: 0.6001
Epoch 4/20
469/469 [=====] - 47s 101ms/step - loss: 0.0925 - acc: 0.6001
Epoch 5/20
469/469 [=====] - 48s 102ms/step - loss: 0.0921 - acc: 0.6001
Epoch 6/20
469/469 [=====] - 48s 103ms/step - loss: 0.0919 - acc: 0.6001
Epoch 7/20
469/469 [=====] - 49s 104ms/step - loss: 0.0919 - acc: 0.6001
Epoch 8/20
469/469 [=====] - 48s 102ms/step - loss: 0.0905 - acc: 0.6001
Epoch 9/20
469/469 [=====] - 47s 100ms/step - loss: 0.0913 - acc: 0.6001
Epoch 10/20
469/469 [=====] - 47s 100ms/step - loss: 0.0900 - acc: 0.6001
Epoch 11/20
469/469 [=====] - 47s 100ms/step - loss: 0.0896 - acc: 0.6001
Epoch 12/20
469/469 [=====] - 47s 100ms/step - loss: 0.0909 - acc: 0.6001
Epoch 13/20
469/469 [=====] - 47s 99ms/step - loss: 0.0893 - acc: 0.6001
Epoch 14/20
469/469 [=====] - 48s 101ms/step - loss: 0.0905 - acc: 0.6001
Epoch 15/20
469/469 [=====] - 48s 102ms/step - loss: 0.0893 - acc: 0.6001
Epoch 16/20
469/469 [=====] - 47s 99ms/step - loss: 0.0897 - acc: 0.6001
Epoch 17/20
469/469 [=====] - 46s 99ms/step - loss: 0.0894 - acc: 0.6001
Epoch 18/20
469/469 [=====] - 46s 99ms/step - loss: 0.0893 - acc: 0.6001
Epoch 19/20
469/469 [=====] - 46s 99ms/step - loss: 0.0886 - acc: 0.6001
Epoch 20/20
469/469 [=====] - 47s 100ms/step - loss: 0.0901 - acc: 0.6001
<tensorflow.python.keras.callbacks.History at 0x7f159fa4a890>

```

```

1 model.save_weights('model.h2')
2 model.load_weights('model.h2')
3 model.evaluate(test_images, test_labels)

```

```

313/313 [=====] - 3s 8ms/step - loss: 0.2068 - acc: 0.6001
[0.2068410962820053, 0.6001999974250793]

```

Visualizing the Tensorboard outcome

```

1 %load_ext tensorboard
2 %tensorboard --logdir log

```

TensorBoard

SCALARS

IMAGES

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting method: **default** ▼

Smoothing



0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs


- ☐ ☐ train
- ☐ ☐ validation

TOGGLE ALL RUNS

log

epoch_acc


epoch_acc



run to download ▼

epoch_loss

epoch_loss



run to download ▼

1

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
Reusing TensorBoard on port 6006 (pid 969), started 0:17:22 ago. (Use '!kill 969

TensorBoard

SCALARS

IMAGES

INACTIVE

Search nodes. Regexes support...

Fit to Screen

Download PNG

Run (1) train

Tag (3) Default

Upload Choose File

Graph

Conceptual Graph

Profile

Trace inputs

Color

Structure

Device

Close legend.

Graph (* = expandable)

Namespace* ?

OpNode ?

Main Graph

1

https://colab.research.google.com/drive/1eC-Nku_kQil6bBTq6OJ1dnY-9C5xP6rC#scrollTo=uHzGbQZh5NUt&printMode=true

8/11

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6006 (pid 969), started 0:17:26 ago. (Use '!kill 969

TensorBoard

SCALARS

IMAGES

INACTIVE

Histogram mode

OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☐ train☐ validation

TOGGLE ALL RUNS

log


Filter tags (regular expressions supported)

conv2d_1

2 ^


conv2d_1/bias_0

train



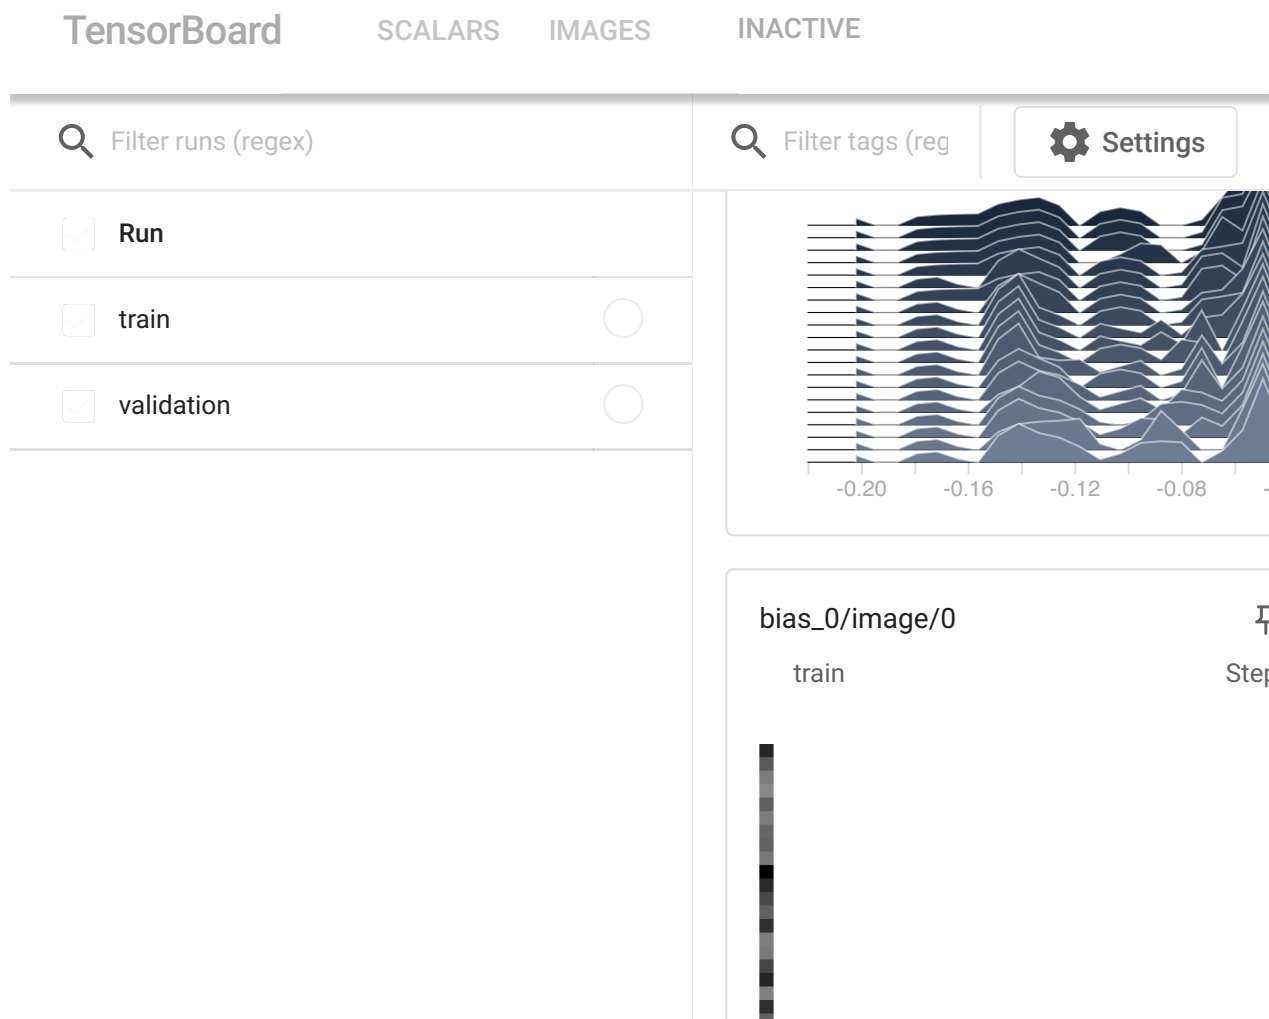
conv2d_1/kernel_0

train



```
%reload_ext tensorboard
```

```
Reusing TensorBoard on port 6006 (pid 969), started 0:23:00 ago. (Use '!kill 969
```



1



